

Faster & Better Power BI

Marco Englaro

Cloud Solution Architect

Microsoft

Who I am

Cloud Solution Architect Data & AI
Co Founder



Marco Englaro



[@MarcoEnglaro](https://twitter.com/MarcoEnglaro)



[Marco Englaro | LinkedIn](https://www.linkedin.com/in/MarcoEnglaro)

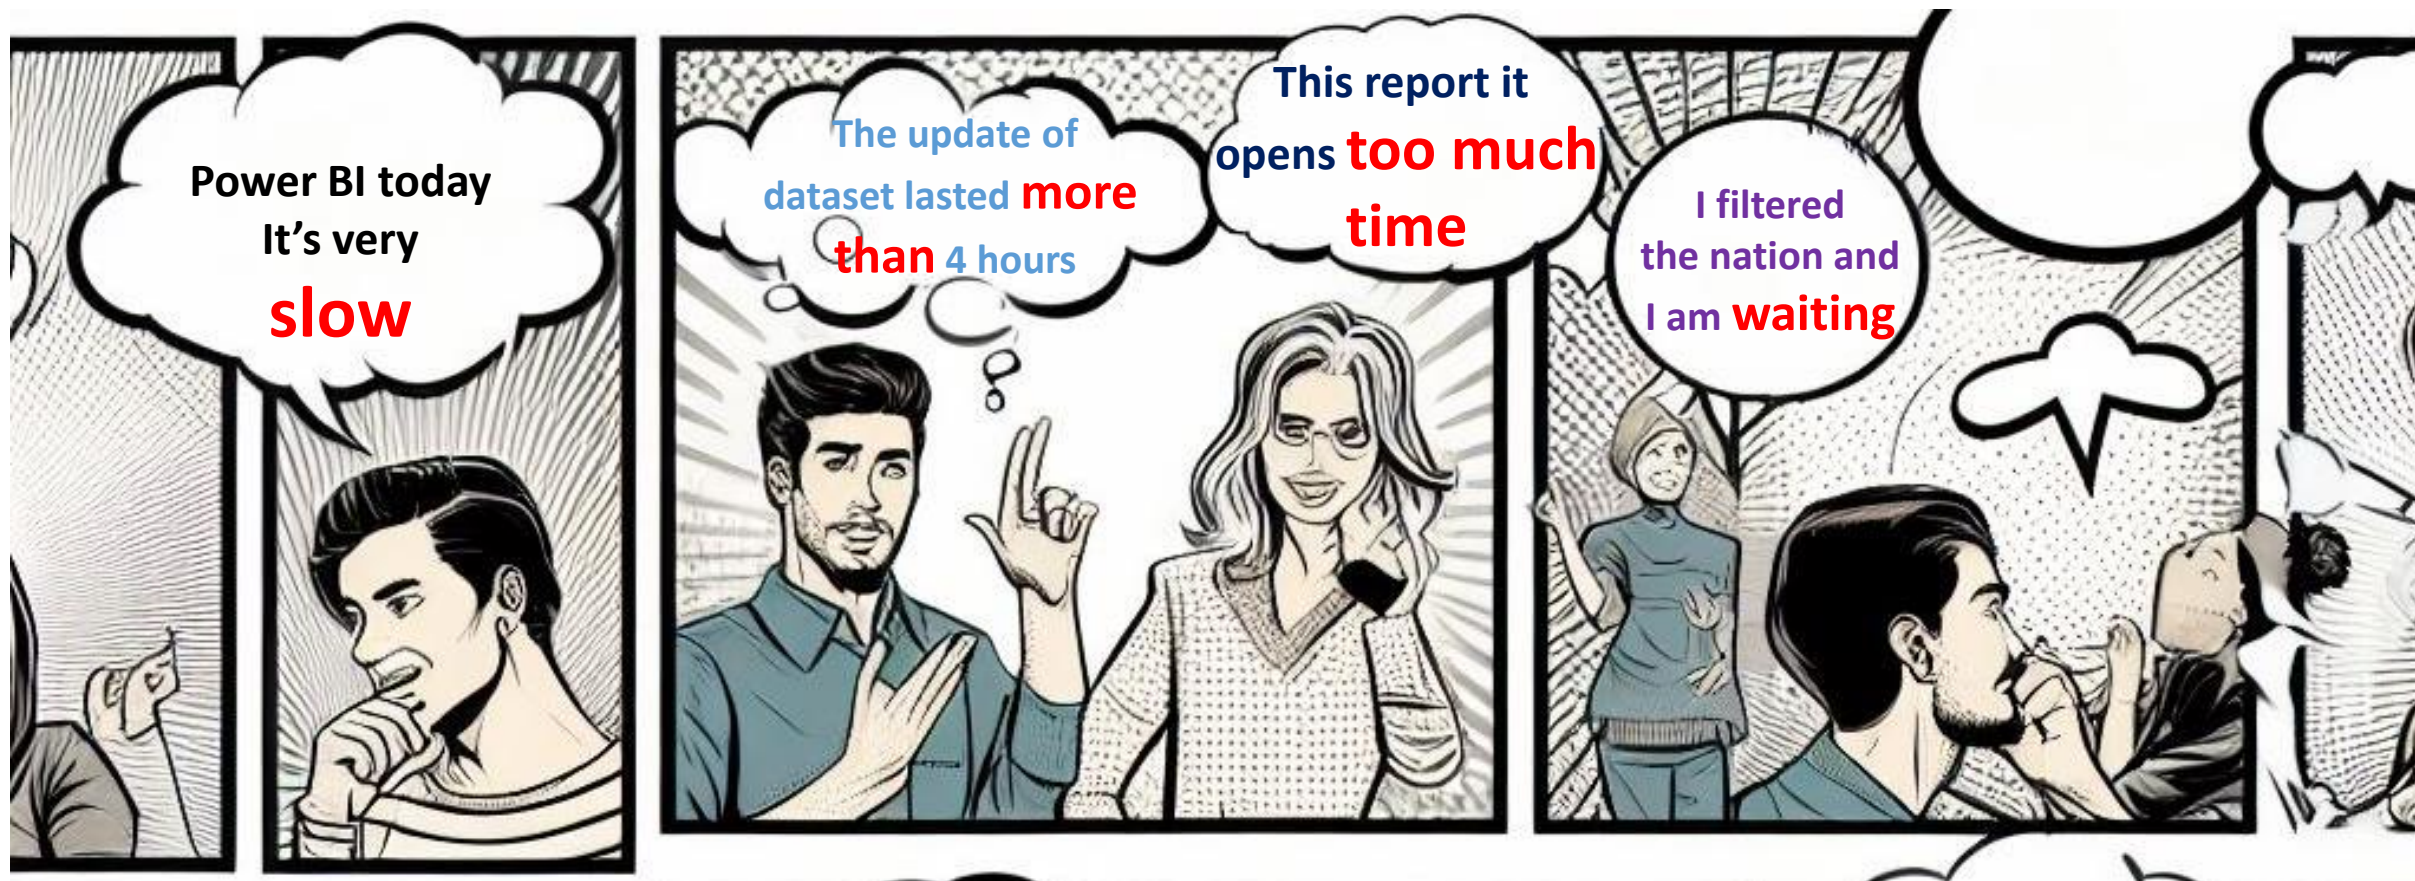
Improve speed performance of Power BI



Power BI



Performance Issues



Right Approach



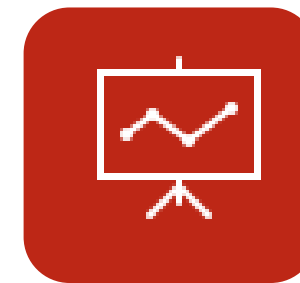
Environment and Architecture

Power BI Capacity



Semantic Model

Connection Type
Data Model



Report

Dax Query
Report Interaction

Environment and Architecture

- Power BI Capacity - Large Dataset

Large semantic model storage format setting **has benefits**.

When enabled, the large semantic model storage format **can improve XMLA write operations performance**.

For semantic models using the large semantic model storage format, Power BI automatically **sets the default segment size to 8 million rows** to strike a good balance between memory requirements and query performance for large tables.

Settings

Adventure Works

About

Premium

Azure connections (preview)

Premium capacity ⓘ

☒ On

Choose an available Premium capacity for this workspace

Power BI - West Central US

Default storage format

Small dataset storage format

Small dataset storage format

Large dataset storage format

powerbi://api.powerbi.com/v1.0/myorg/Adventure Works

Copy

Delete workspace

Save

Cancel



Environment and Architecture

- Power BI Capacity - Autoscale

Automatically use **more v-cores** (virtual CPU cores) when the computing load on your Power BI Premium subscription would **otherwise be slowed by its capacity.**



The screenshot displays the Power BI Admin portal interface. On the left is a navigation sidebar with options like Tenant settings, Usage metrics, Users, Premium Per User (preview), Audit logs, Capacity settings (selected), Refresh summary, Embed Codes, Organizational visuals, Azure connections (preview), Workspaces, Custom branding, Protection metrics, and Featured content. The main content area is titled 'Power BI Premium > proffysworldbyoladxt'. It features a 'Premium Generation 2 (preview)' section with an 'Enabled' toggle. Below this are two large cards: 'Size | P1' showing '8 Base v-cores' with a 'Change size' button, and 'Auto scale | Off' showing '0 Additional v-cores in use' with a 'Manage auto-scale' button. Explanatory text and links are provided for both settings. At the bottom, there is a link to 'Capacity usage report'.

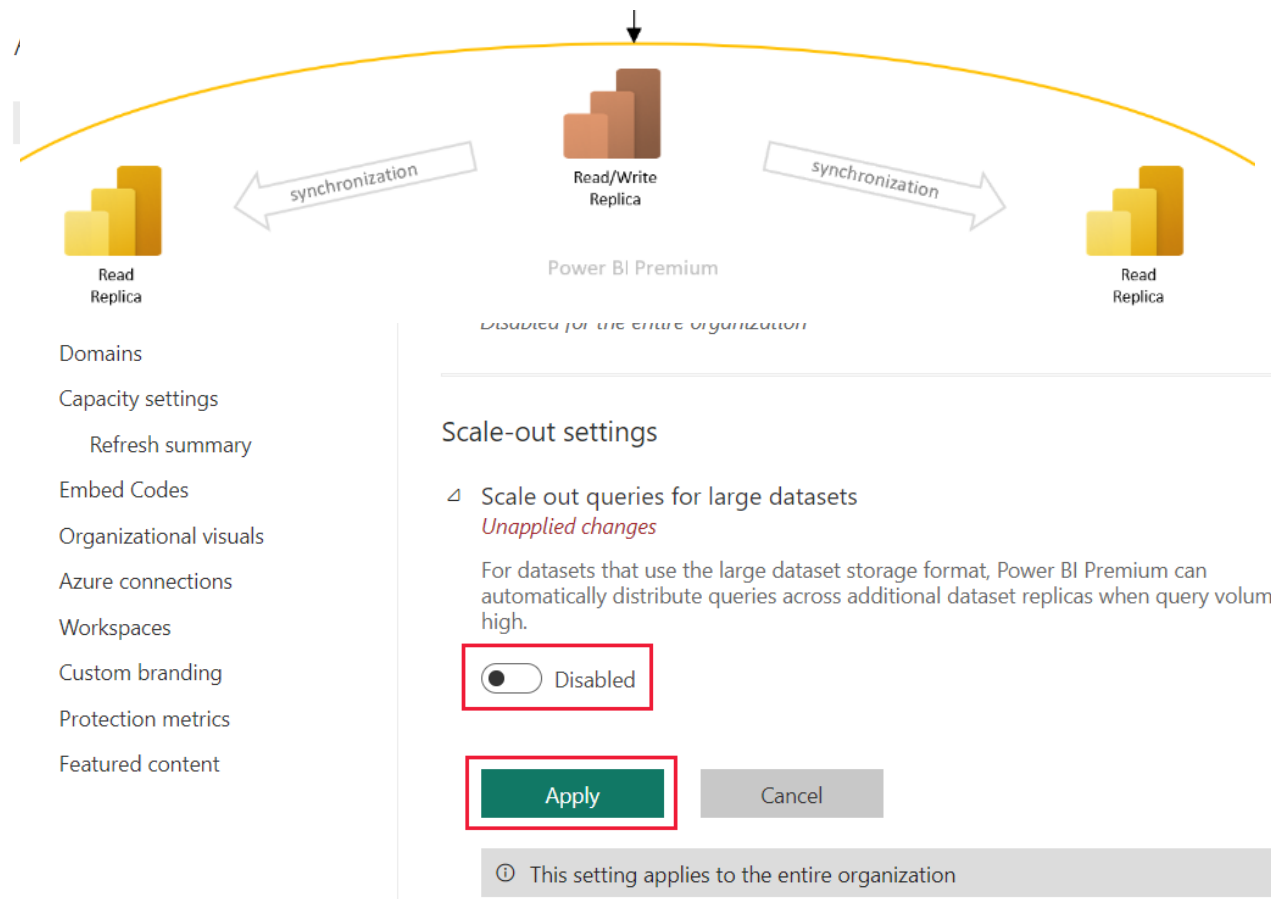
Environment and Architecture

- Power BI Capacity - **Scale Out**

Semantic model scale-out **helps Power BI deliver fast performance** while your reports and dashboards are consumed **by a large audience**.

Semantic model scale-out uses your Premium capacity to **host one or more read-only replicas of your primary semantic model**.

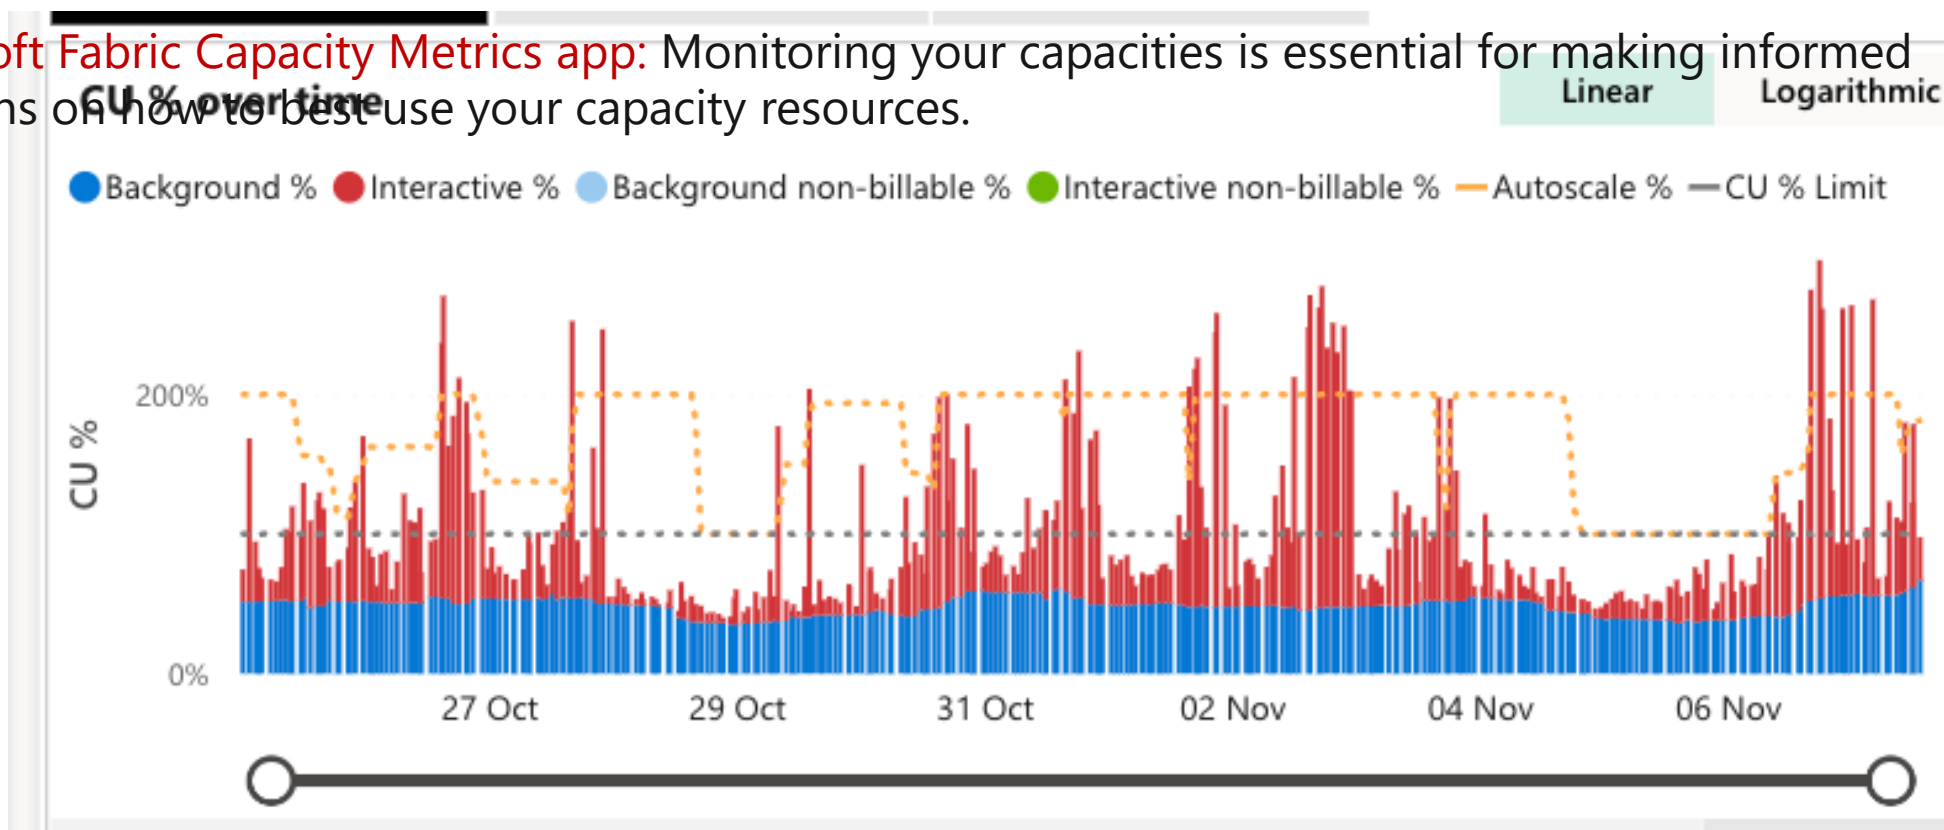
By increasing throughput, the read-only replicas ensure performance **doesn't slow down** when multiple users submit queries at the same time.



Environment and Architecture

- Power BI Capacity - Monitoring

Microsoft Fabric Capacity Metrics app: Monitoring your capacities is essential for making informed decisions on how to best use your capacity resources.



Right Approach



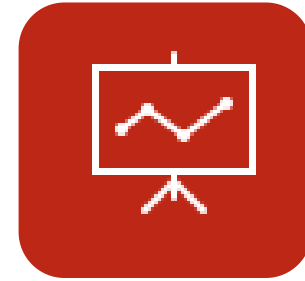
Environment and Architecture

Power BI Capacity
Gateway Configuration
Driver Configuration



Semantic Model

Connection Type
Data Model



Report

Dax Query
Report Interaction

Environment and Architecture



- Gateway Configuration - Connection type

Separating sources prevents the gateway from having thousands of DirectQuery requests queued up at the same time as the morning's scheduled refresh of a large-size data model that's used for the company's main dashboard.

Schedule Refresh

Depending on your query size and the number of refreshes that occur per day, you can choose to stay with the recommended minimum hardware requirements or upgrade to a higher performance machine. If a given query isn't folded, transformations occur on the gateway machine. As a result, **the gateway machine benefits from having more available RAM.**

Direct Query

A query is sent each time any user opens the report or looks at data. If you expect more than 1,000 users to access the data concurrently, make sure your computer has robust and capable hardware components. **More CPU cores result in better throughput for a DirectQuery connection.**

Environment and Architecture

2023 WPC



- Gateway Configuration - Location

The location of the gateway installation can have **significant effect on your query performance**.

Try to make sure that your gateway, data source locations, and the Power BI tenant are as close as possible to each other to **minimize network latency**.

To determine your Power BI tenant location, in the Power BI service select the question mark (?) icon in the upper-right corner. Then select **About Power BI**.



Power BI

©Microsoft Corporation 2023. All rights reserved.

Service version: 13.0.22007.81

Client version: 2311.1.16680-train

Activity ID: fec37f4-57a2-4a11-ba3f-3e8ac4a89634

App Instance ID: wcd6v

User object ID: 3a453d62-107e-4e95-962b-f6c96a187e46

Tenant URL: <https://app.powerbi.com/home?ctid=f6eeb23a-b8bd-4>

Your data is stored in West Europe (Netherlands)

Thu Nov 23 2023 23:49:30 GMT+0100 (Ora standard dell'Europa centrale)

Environment and Architecture

2023 WPC



- Gateway Configuration - Monitoring



Rui Romano
Principal Program Manager



Environment and Architecture



- Driver Configuration

Managed
Driver

Microsoft
Documentation

Vendor
Documentation

Avoid Generic
Drivers

Right Approach



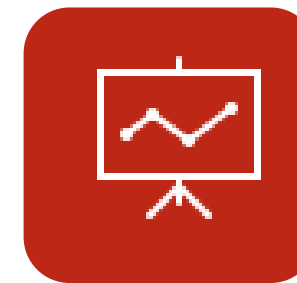
Environment and Architecture

Power BI Capacity
Gateway Configuration
Driver Configuration



Semantic Model

Connection Type
Data Model



Report

Dax Query
Report Interaction

Semantic Model – Connecton Type

- not everything is the same

Import
Mode

Direct Query
Mode

Composite
Mode

Mixed Mode

Direct Lake

Semantic Model – Data Model

1. Import only those columns you really need
2. Reduce the column cardinality!
3. Same as for columns, keep only those rows you need
4. Avoid using calculated columns whenever possible, since they are not being optimally compressed
5. Use proper data types (DateTime to Date)
6. Avoid using floating point data types
7. Disable Auto Date/Time option for data loading
8. Use Star schema instead of Snowflake schema when possible
9. Avoid bi-directional and many-to-many relationships against high cardinality columns
10. Overusing Calculated Columns
11. Overusing expensive relationships
12. Remove Primary Key columns from fact tables

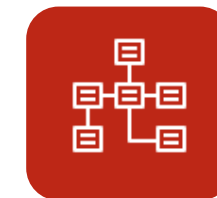
Semantic Model – Data Model

Reduce the column cardinality!

High cardinality columns in PBI models can be particularly expensive. The best practice is to remove them from the model, especially when these columns are not relevant for data analysis, such as a **GUID or timestamp** of a SQL Server table.

However, whenever the information they contain is required, you can optimize these columns **by splitting the value in two or more columns** with a smaller number of distinct values.

This will require some more effort when accessing the column value, but the saving can be so high in large tables that it **could definitely worth the effort**.



Semantic Model – Data Model

Avoid using calculated columns whenever possible, since they are not being optimally compressed

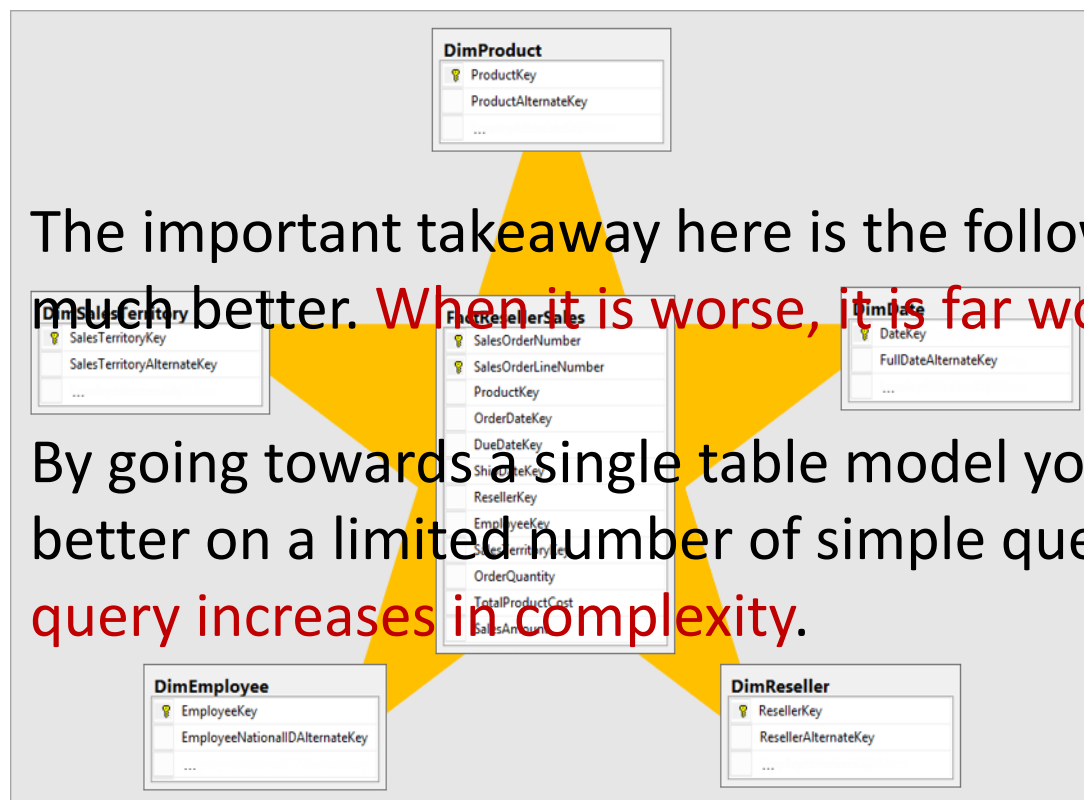
Created at
the Data
Source

Created
with
Power
Query

Created in
DAX

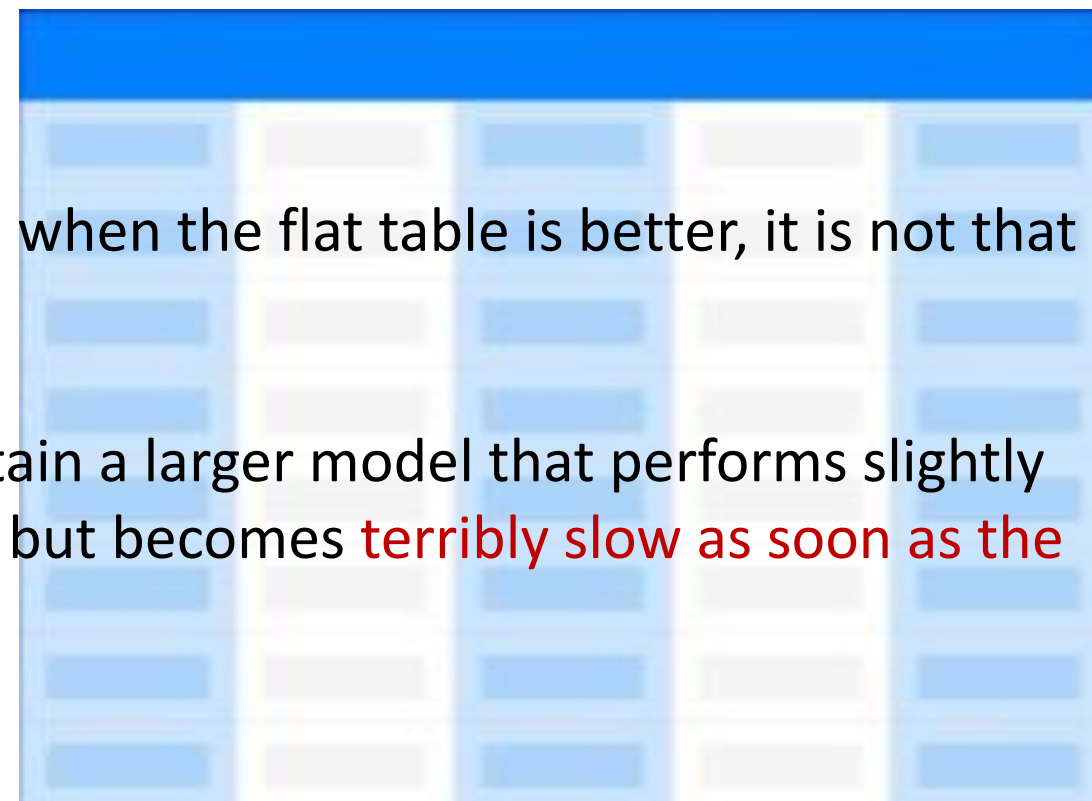
Semantic Model – Data Model

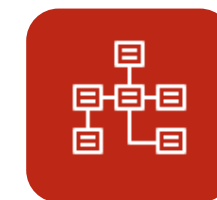
Use Star schema instead of flat table



The important takeaway here is the following: when the flat table is better, it is not that much better. **When it is worse, it is far worse!**

By going towards a single table model you obtain a larger model that performs slightly better on a limited number of simple queries, but becomes **terribly slow as soon as the query increases in complexity.**



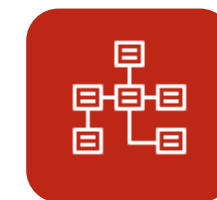


Semantic Model – Data Model

External Tools



Best Practice Analyzer	
Show ignored	
Object	Type
+ [Performance] Set IsAvailableInMdx to false on non-attribute columns (6 objects)	
+ [Performance] Model should have a date table (1 object)	
+ [Performance] Date/calendar tables should be marked as a date table (1 object)	
+ [Performance] Minimize Power Query transformations (1 object)	
+ [Performance] Reduce usage of calculated tables (1 object)	
+ [Maintenance] Remove unnecessary columns (3 objects)	
+ [Maintenance] Ensure tables have relationships (1 object)	
+ [Maintenance] Visible objects with no description (31 objects)	
+ [Naming Conventions] Partition name should match table name for single partition tables (2 objects)	
+ [Formatting] Provide format string for "Date" columns (4 objects)	
+ [Formatting] Do not summarize numeric columns (1 object)	
+ [Formatting] Whole numbers should be formatted with thousands separators and no decimals (6 objects)	
+ [Formatting] Relationship columns should be of integer data type (2 objects)	
+ [Formatting] Add data category for columns (2 objects)	
+ [Formatting] Hide foreign keys (3 objects)	
+ [Formatting] Mark primary keys (3 objects)	
+ [Formatting] Hide fact table columns (1 object)	
69 objects in violation of 17 Best Practice rules.	



Semantic Model – Data Model

External Tools



Log Results History ● Server Timings VertiPaq Analyzer

	Name	Cardinality	Total Size ↓	Data	Dictionary	Hier Size	Encoding	Data Type	RI Violations	User Hier Size	Rel Size	% Table	% DB	
Columns	▶ Sales	225.616.948	2.624.526.272	2.561.906.9...	42.341.056	15.287.296	Many	-	-	0	4.990.936		92,79%	
Relationships	▶ Customer	1.868.084	201.471.877	26.116.400	137.940.517	37.414.960	Many	-	-	0	0		7,12%	
	▶ CurrencyExchan...	202.100	1.281.712	728.352	396.152	144.264	Many	-	-	0	12.944		0,05%	
Partitions	▶ Product	2.517	510.992	40.912	403.728	66.352	Many	-	-	0	0		0,02%	
Summary	▶ Date	4.018	423.066	41.848	333.346	47.872	Many	-	1	0	0		0,01%	
	▶ Store	74	82.544	2.024	78.064	2.456	Many	-	-	0	0		0,00%	
	▶ Date DQ	0	10.504	2.312	8.192	0	DQ	-	-	0	0		0,00%	
	▶ Product DQ	0	9.992	2.040	7.952	0	DQ	-	-	0	0		0,00%	
	▶ Customer DQ	0	9.488	1.904	7.584	0	DQ	-	-	0	0		0,00%	
	▶ Sales DQ	0	7.184	1.904	5.256	0	DQ	-	-	0	24		0,00%	
	▶ CurrencyExchan...	0	2.928	680	2.240	0	DQ	-	-	0	8		0,00%	

Right Approach



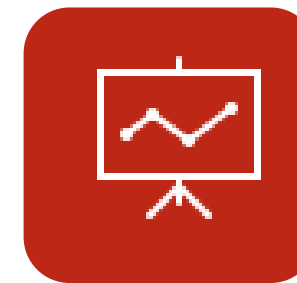
Environment and Architecture

Power BI Capacity
Gateway Configuration
Driver Configuration



Semantic Model

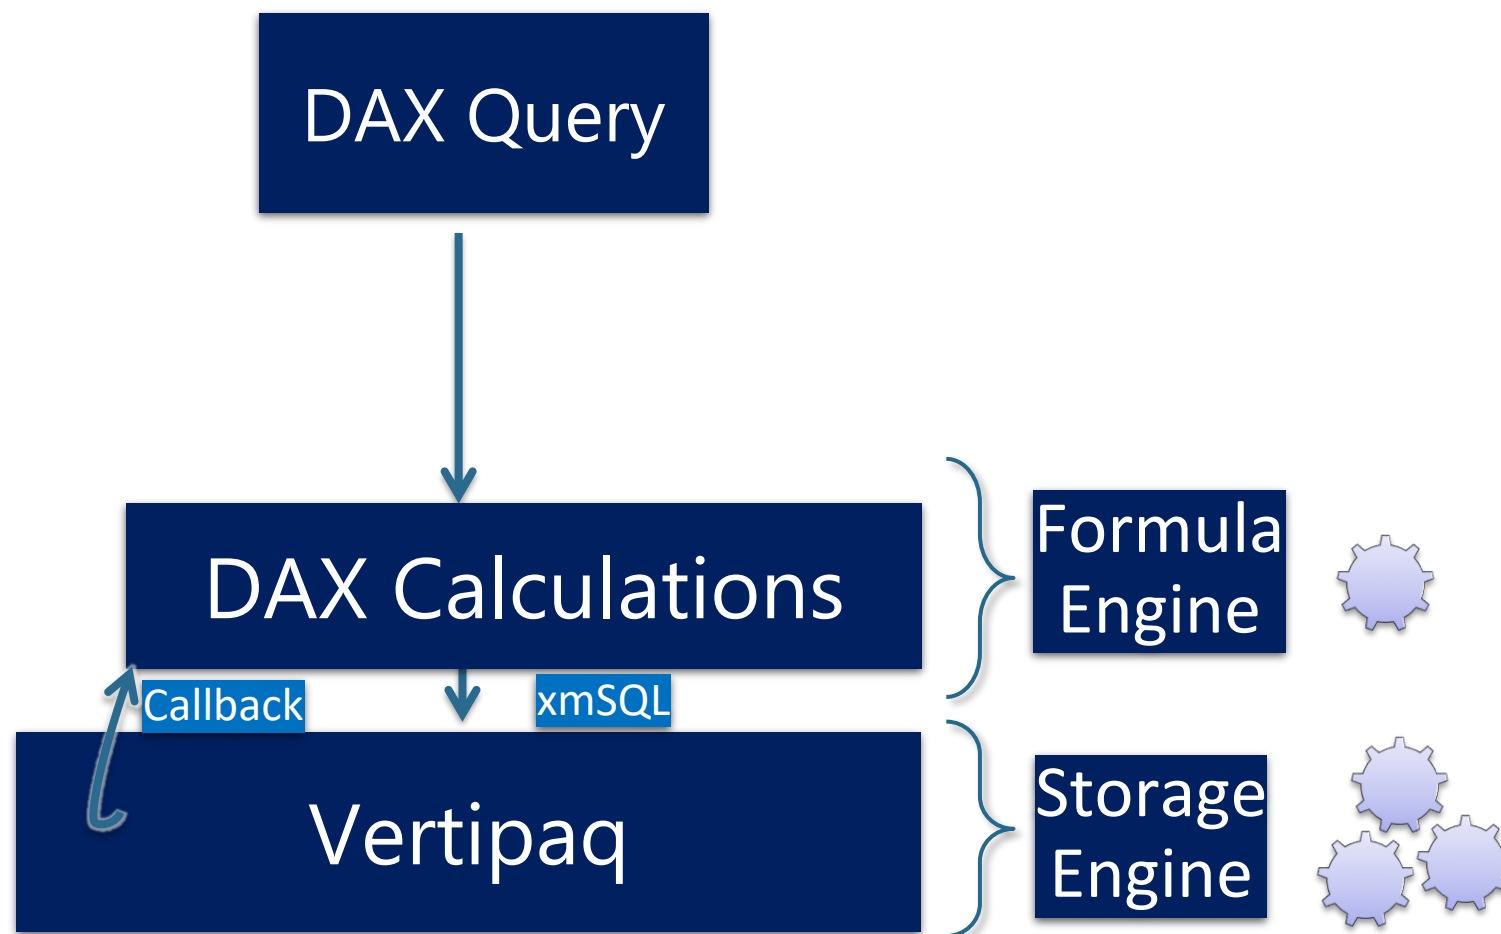
Connection Type
Data Model



Report

Dax Query
Report Interaction

Report – Query Architecture



FORMULA ENGINE (FE)

- Monothreaded, ask and fetch SE queries sequentially
- Very Smart !
- Build execution plans
- No cache at FE level for DAX queries

STORAGE ENGINE (SE)

- Multithreaded, scan segments in parallel
- Very Fast !
- Can only handle simple arithmetic calculations (callback can be needed)
- SE query results are cached (if no callback)

Report – Dax Optimize

Sales Table: 225,616,948 Milion Rows

Year	Sales Gt 200
2010	5.966.043.126,80
Aug	676.336.337,95
Dec	1.353.810.935,74
Jul	638.715.185,11
Jun	666.241.293,46
May	303.911.689,67
Nov	782.639.825,22
Oct	785.281.421,80
Sep	759.106.437,86
2011	10.691.471.516,63
2012	12.627.251.409,30
2013	20.443.916.137,00
2014	29.136.505.467,53
2015	23.549.371.252,73
2016	17.946.710.230,43
2017	30.342.774.851,17
2018	38.939.706.798,16
2019	27.577.227.619,71
2020	6.038.286.952,60
Total	223.259.265.362,06

It all starts with a report that users describe as slow. It contains a simple measure that computes sales for only transactions whose amount is greater than 200 US

DAX MEASURE

```

Sales Gt 200 = SUMX (Sales, IF (Sales[Quantity] * Sales[Net
Price] >= 200, Sales[Quantity] * Sales[Net Price]))
Sales,
IF (
    Sales[Quantity] * Sales[Net Price] >= 200,
    Sales[Quantity] * Sales[Net Price]
)
)

```



Report – Dax Optimize

Sales Table: 225,616,948 Milion Rows

Year	Sales Gt 200
2010	5.966.043.126,80
Aug	676.336.337,95
Dec	1.353.810.935,74
Jul	638.715.185,11
Jun	666.241.293,46
May	303.911.689,67
Nov	782.639.825,22
Oct	785.281.421,80
Sep	759.106.437,86
2011	10.691.471.516,63
2012	12.627.251.409,30
2013	20.443.916.137,00
2014	29.136.505.467,53
2015	23.549.371.252,73
2016	17.946.710.230,43
2017	30.342.774.851,17
2018	38.939.706.798,16
2019	27.577.227.619,71
2020	6.038.286.952,60
Total	223.259.265.362,06

We use the **Performance Analyzer tool** in Power BI Desktop to retrieve the DAX query executed for the visual. The query took 6 seconds to run.

Performance analyzer	
Start recording	Refresh visuals Stop
Clear Export	
Name	Duration (ms)
Recording started (23/11/2023 10:51:11)	-
Refreshed visual	-
Matrix	6175
DAX query	6068
Visual display	32
Other	75
Copy query	

Report – Dax Optimize

Sales Table: 225,616,948 Milion Rows

Year	Sales Gt 200
2010	5.966.043.126,80
Aug	676.336.337,95
Dec	1.353.810.935,74
Jul	638.715.185,11
Jun	666.241.293,46
May	303.911.689,67
Nov	782.639.825,22
Oct	785.281.421,80
Sep	759.106.437,86
2011	10.691.471.516,63
2012	12.627.251.409,30
2013	20.443.916.137,00
2014	29.136.505.467,53
2015	23.549.371.252,73
2016	17.946.710.230,43
2017	30.342.774.851,17
2018	38.939.706.798,16
2019	27.577.227.619,71
2020	6.038.286.952,60
Total	223.259.265.362,06

The first thing to do is to execute the query in DAX Studio with Server Timings enabled to obtain the first baseline. Later on, we will check that the simplified query did not change the timings in such a way that the issue seems resolved. Here is the DAX Studio timings report

Log Results History **Server Timings**

Record

Pause

Stop

Clear

Copy

Export

Info

Total 8.907 ms	SE CPU 56.110 ms x6,3
<div><div></div><div>FE 9 ms 0,1%</div></div>	<div><div></div><div>SE 8.898 ms 99,9%</div></div>
<div></div>	
SE Queries 3	SE Cache 0 0,0%

Line	Subclass	Duration	CPU	Par.	Rows	KB	Timeline
2	Scan	3.057	19.172	x6,3	3.150	50	<div></div>
4	Scan	3.001	18.750	x6,2	14	1	<div></div>
6	Scan	2.840	18.188	x6,4	1	1	<div></div>



Report – Dax Optimize

Sales Table: 225,616,948 Milion Rows

Year Sales Gt 200

2010	5.966.043.126,80
Aug	676.336.337,95
Dec	1.353.810.935,74
Jul	638.715.185,11
Jun	666.241.293,46
May	303.911.689,67
Nov	782.639.825,22
Oct	785.281.421,80
Sep	759.106.437,86
2011	10.691.471.516,63
2012	12.627.251.409,30
2013	20.443.916.137,00
2014	29.136.505.467,53
2015	23.549.371.252,73
2016	17.946.710.230,43
2017	30.342.774.851,17
2018	38.939.706.798,16
2019	27.577.227.619,71
2020	6.038.286.952,60
Total	223.259.265.362,06

Log Results History **Server Timings**

Record

Pause

Stop

Clear

Copy

Export

Info

Total
8.907 ms

SE CPU
56.110 ms
x6,3

FE
9 ms
0,1%

SE
8.898 ms
99,9%

SE Queries
3

SE Cache
0
0,0%

Line	Subclass	Duration	CPU	Par.	Rows	KB	Timeline	Query
2	Scan	3.057	19.172	x6,3	3.150	50	<div></div>	WITH \$Expr0 := [C
4	Scan	3.001	18.750	x6,2	14	1	<div></div>	WITH \$Expr0 := [C
6	Scan	2.840	18.188	x6,4	1	1	<div></div>	WITH \$Expr0 := [C

SET DC_KIND="AUTO";

WITH

\$Expr0 := [CallbackDataID (IF (

Sales[Quantity] * Sales[Net Price] >= 200,

Sales[Quantity] * Sales[Net Price]

))] (PFDATAID ('Sales'[Quantity]) , PFDATAID ('Sales'[Net Price]))

SELECT

'Date'[Year],

'Date'[Month],

'Date'[Month Number],

Though the entire execution time is reported as storage engine CPU, we can clearly see a CallbackDataID, indicating that the formula engine is required to kick in to compute expressions that cannot be pushed down to the storage engine.

Report – Dax Optimize

Sales Table: 225,616,948 Million Rows

Year Sales Gt 200

2010	5.966.043.126,80
Aug	676.336.337,95
Dec	1.353.810.935,74
Jul	638.715.185,11
Jun	666.241.293,46
May	303.911.689,67
Nov	782.639.825,22
Oct	785.281.421,80
Sep	759.106.437,86
2011	10.691.471.516,63
2012	12.627.251.409,30
2013	20.443.916.137,00
2014	29.136.505.467,53
2015	23.549.371.252,73
2016	17.946.710.230,43
2017	30.342.774.851,17
2018	38.939.706.798,16
2019	27.577.227.619,71
2020	6.038.286.952,60
Total	223.259.265.362,06

The problem is the IF statement inside the iteration carried on by SUMX, because the VertiPaq storage engine does not support conditional logic. We must rephrase the measure to avoid the IF statement; we replace it with a condition set by CALCULATE to rely on filtering rather than IF.

Sales Gt 200 =

```

SUMX(
    CALCULATE (
        Sales,
        IF (FILTER ( Sales, Sales[Quantity] * Sales[Net Price] >= 200 )
            Sales[Quantity] * Sales[Net Price] >= 200,
            Sales[Quantity] * Sales[Net Price]
        )
    )
)

```

Report – Dax Optimize

Sales Table: 225,616,948 Milion Rows

Year	Sales Gt 200
2010	5.966.043.126,80
Aug	676.336.337,95
Dec	1.353.810.935,74
Jul	638.715.185,11
Jun	666.241.293,46
May	303.911.689,67
Nov	782.639.825,22
Oct	785.281.421,80
Sep	759.106.437,86
2011	10.691.471.516,63
2012	12.627.251.409,30
2013	20.443.916.137,00
2014	29.136.505.467,53
2015	23.549.371.252,73
2016	17.946.710.230,43
2017	30.342.774.851,17
2018	38.939.706.798,16
2019	27.577.227.619,71
2020	6.038.286.952,60
Total	223.259.265.362,06

The storage engine CPU is a bit lower than the previous version of the measure, but the degree of parallelism is much lower this time. Moreover, the formula engine executes a significant portion of code, making the overall performance much worse than the previous one. Overall, the execution time went from 8 to 29 seconds

Log Results History **Server Timings**

Record

Pause

Stop

Clear

Copy

Export

Info

Total

29.633 ms

SE CPU

69.766 ms

x4,4

FE

13.647 ms

46,1%

SE

15.986 ms

53,9%

SE Queries

5

SE Cache

0

0,0%

Line

Subclass

Duration

CPU

Par.

Rows

KB

Timeline

Query

2

Scan

5.248

18.453

x3,5

5.150.659

60.360

WITH

4

Scan

4.976

18.891

x3,8

5.150.659

20.120

SELE

6

Scan

2.930

16.078

x5,5

735.448

8.619

WITH

8

Scan

2.124

12.188

x5,7

735.448

2.873

SELE

10

Scan

708

4.156

x5,9

1

1

WITH

Report – Dax Optimize

Sales Table: 225,616,948 Milion Rows

Year	Sales Gt 200
2010	5.966.043.126,80
Aug	676.336.337,95
Dec	1.353.810.935,74
Jul	638.715.185,11
Jun	666.241.293,46
May	303.911.689,67
Nov	782.639.825,22
Oct	785.281.421,80
Sep	759.106.437,86
2011	10.691.471.516,63
2012	12.627.251.409,30
2013	20.443.916.137,00
2014	29.136.505.467,53
2015	23.549.371.252,73
2016	17.946.710.230,43
2017	30.342.774.851,17
2018	38.939.706.798,16
2019	27.577.227.619,71
2020	6.038.286.952,60
Total	223.259.265.362,06

Indeed, we used a filter over Sales as a filter argument in CALCULATE. A table filter is a very bad practice that newbies oftentimes use. A filter in CALCULATE should work on the minimum number of columns required to obtain its effect

Sales Gt 200 =

```

CALCULATE (
    SUMX ( Sales, Sales[Quantity] * Sales[Net Price] ),
    FILTER ( Sales, Sales[Quantity] * Sales[Net Price] >= 200 )
) ALL (Sales[Quantity] , Sales[Net Price]),
Sales[Quantity] * Sales[Net Price] >= 200 )
)

```

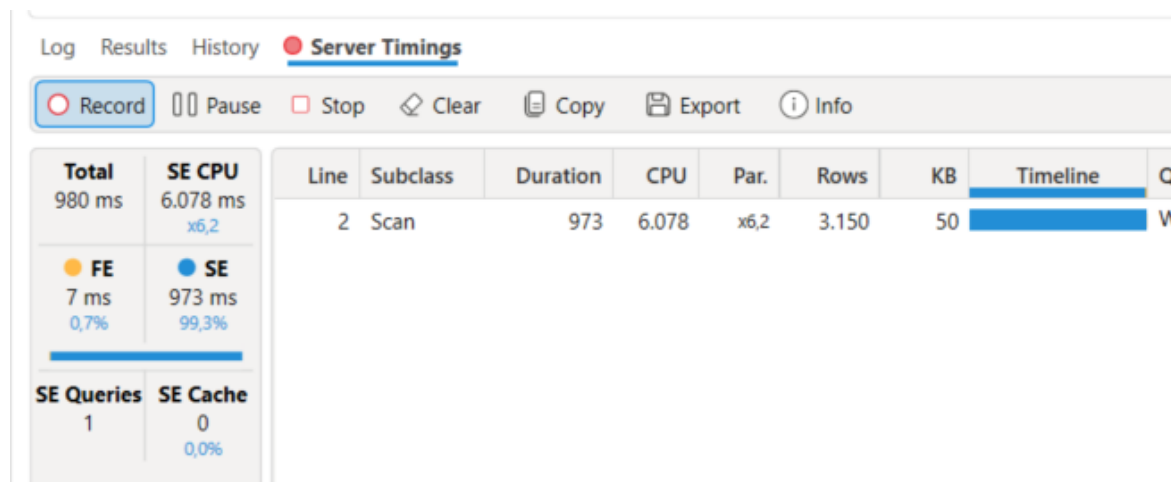
Report – Dax Optimize

Sales Table: 225,616,948 Milion Rows

Year	Sales Gt 200
2010	5.966.043.126,80
Aug	676.336.337,95
Dec	1.353.810.935,74
Jul	638.715.185,11
Jun	666.241.293,46
May	303.911.689,67
Nov	782.639.825,22
Oct	785.281.421,80
Sep	759.106.437,86
2011	10.691.471.516,63
2012	12.627.251.409,30
2013	20.443.916.137,00
2014	29.136.505.467,53
2015	23.549.371.252,73
2016	17.946.710.230,43
2017	30.342.774.851,17
2018	38.939.706.798,16
2019	27.577.227.619,71
2020	6.038.286.952,60
Total	223.259.265.362,06

All the indicators are just perfect. The storage engine CPU is massively reduced, the degree of parallelism is back to being exceptional, there is virtually no formula engine involved in the query and no CallbackDataIDs anywhere.

Materialization is reduced from 5 million rows to only 3,150 rows



Report – Dax Query

1. Use Variables when you can precompute and reuse some calculations
2. Try to avoid CallBack (ex : complex filter)
3. Avoided this function SEARCH, IFERROR, CONTAINS, and INTERSECT functions
4. The FILTER function is often overused. Its main purpose is for filtering columns based on measure values.
5. Add column and measure references in your DAX expressions
6. Use ISBLANK() instead of =Blank() check
7. Use COUNTROWS instead of COUNT
8. Work upstream, if possible
9. format your code! (Use DAX Formatter)
10. Split your calculations in smaller blocks

Performance Issues



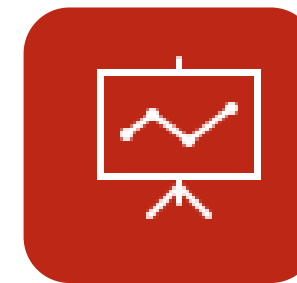
Environment and Architecture

Power BI Capacity
Gateway Configuration
Driver Configuration



Semantic Model

Connection Type
Data Model



Report

Dax Query
Report Interaction

Report – Report Interaction

The problem behind this report being slow is not to be found anywhere in the DAX code or in the model. This report is slow because there are too many card visuals. In that scenario, the solution would be to use small multiples rather than focusing on the DAX code of the measures.

Sales Dashboard

Red products

0,00	20,45K	44,50K	8,46K	14,23K	876,65
6,04K	3,79K	325,71	0,00	2,78K	

Black products

22,75K	80,95K	95,83K	82,16K	24,36K	94,99
73,49K	81,42K	5,24K	37,71K	94,08K	

Blue products

1,93K	9,12K	17,64K	63,97K	14,78K	64,66K
4,63K	1,42K	7,77K	0,00	27,42K	

Year

- ☐ 2007
☐ 2008
☐ 2009
☐ 2010

Report – Report Interaction

Sales Dashboard

Red products

0,00	20,45K
6,04K	3,79K

Black products

22,75K	80,95K
73,49K	81,42K

Blue products

1,93K	9,12K
4,63K	1,42K

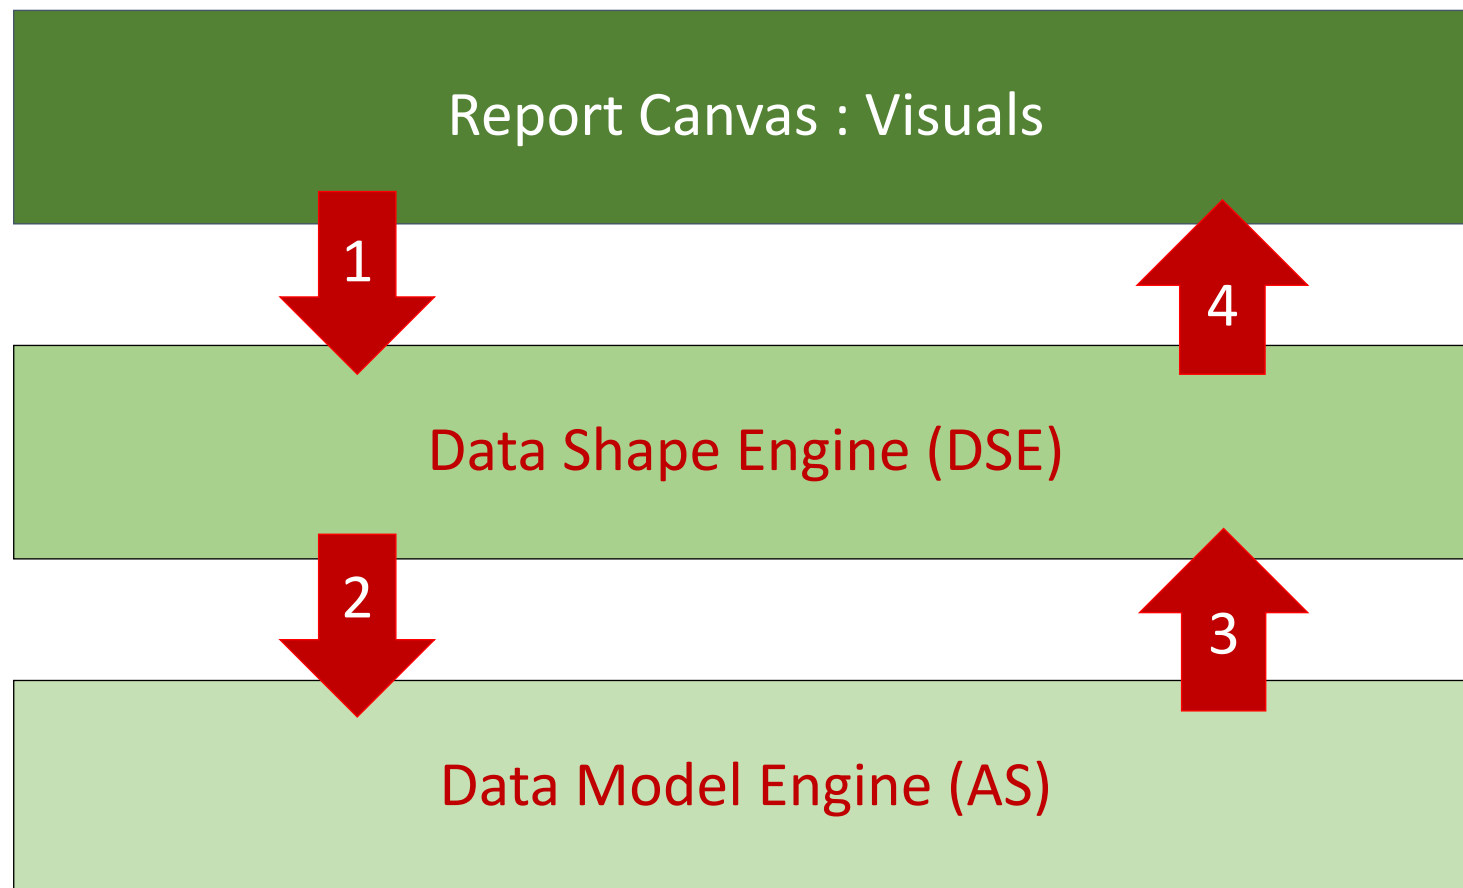
Performance analyzer	
Start recording Refresh visuals Stop	
Clear Export	
Name	Duration (ms)
Red - World Wide Importers - Card	596
Text box	201
Visual display	162
Other	39
Copy query	
Run in DAX Query View	
Black - A. Datum - Card	597
DAX query	67
Visual display	56
Other	474
Copy query	
Run in DAX Query View	
Black - Adventure Works - Card	598
Black - Contoso - Card	599
Black - Fabrikam - Card	600

DAX query: this is the time required to execute the DAX query.

Visual display: indicates the time required to render the visual. On more complex visuals, the rendering time might be significant.

Other: this is the time the visual had to wait before Power BI Desktop could execute the DAX query to populate the visual

Report – Report Interaction



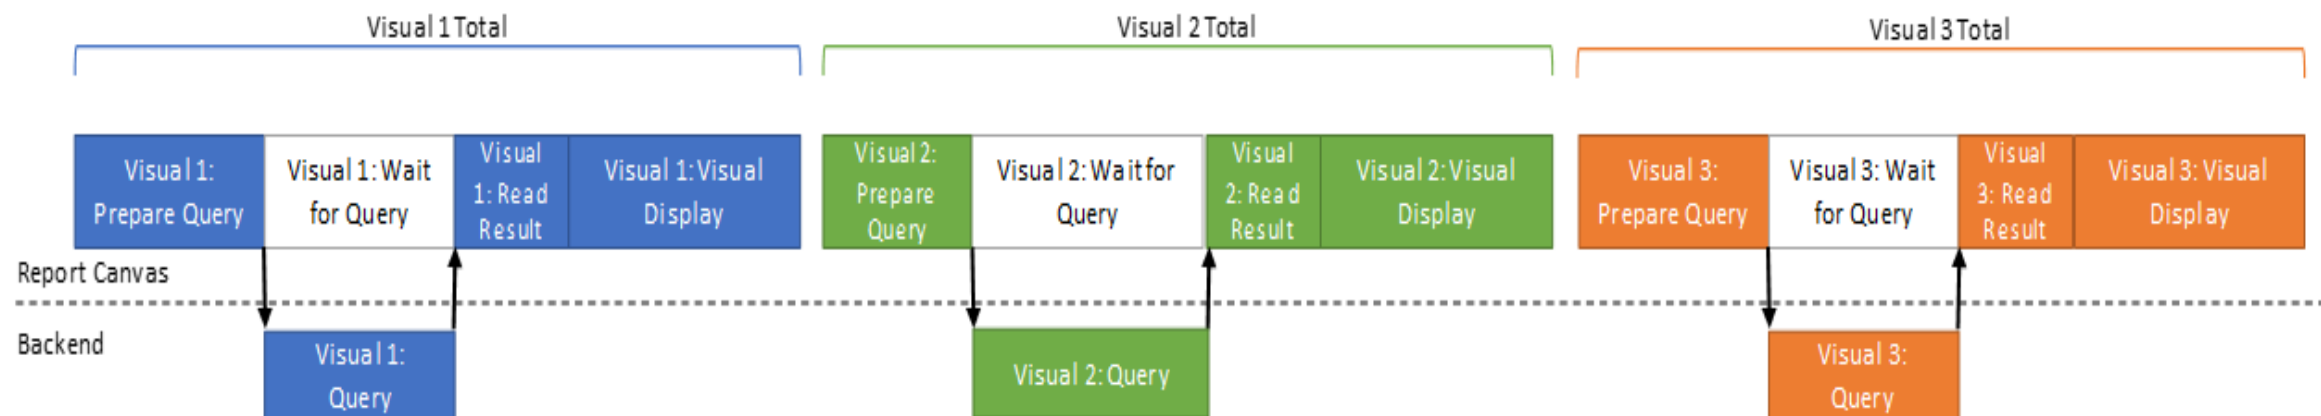
Execution Workflow

Report Canvas retrieves data using an internal, high-level Power BI query language known as Semantic Query.

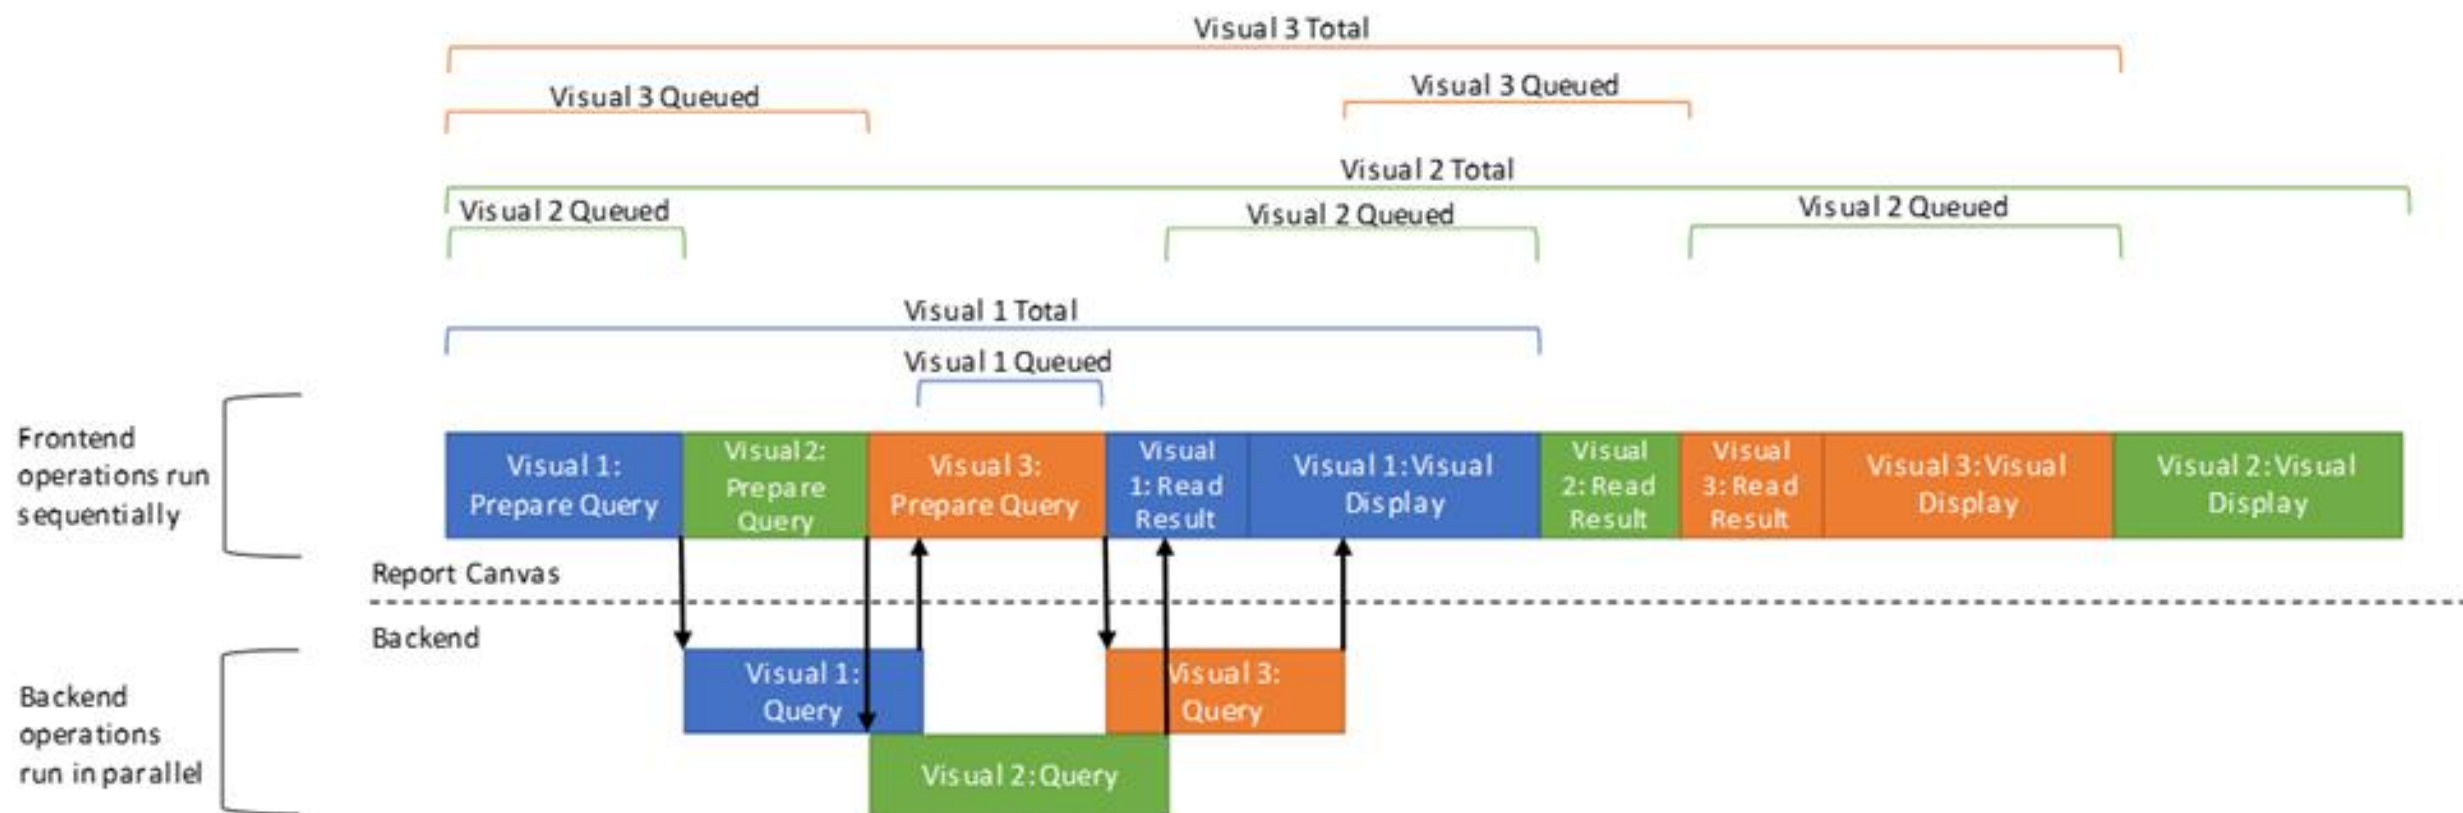
Data Shape Engine (DSE) evaluates semantic queries by generating and executing one or more DAX queries against a semantic model.

Data Model Engine (AS) stores the data model and provides reporting services, such as DAX query evaluation.

Report – Report Interaction



Report – Report Interaction



Report – Report Interaction

Sales Dashboard

Red products

0.00	20.45K	44.50K	8.46K	14.23K	876.64
6.04K	3.79K	325.71	0.00	2.78K	

Black products

22.75K	80.95K	95.83K	82.16K	24.36K	94.99
73.49K	81.42K	5.24K	37.71K	94.08K	

Blue products

1.93K	9.12K	17.64K	63.97K	14.78K	64.66K
4.63K	1.42K	7.77K	0.00	27.42K	

sqlbi

Year

☐ 2007

☐ 2008

☐ 2009

☐ 2010

Performance analyzer

Start recording Refresh visuals Stop

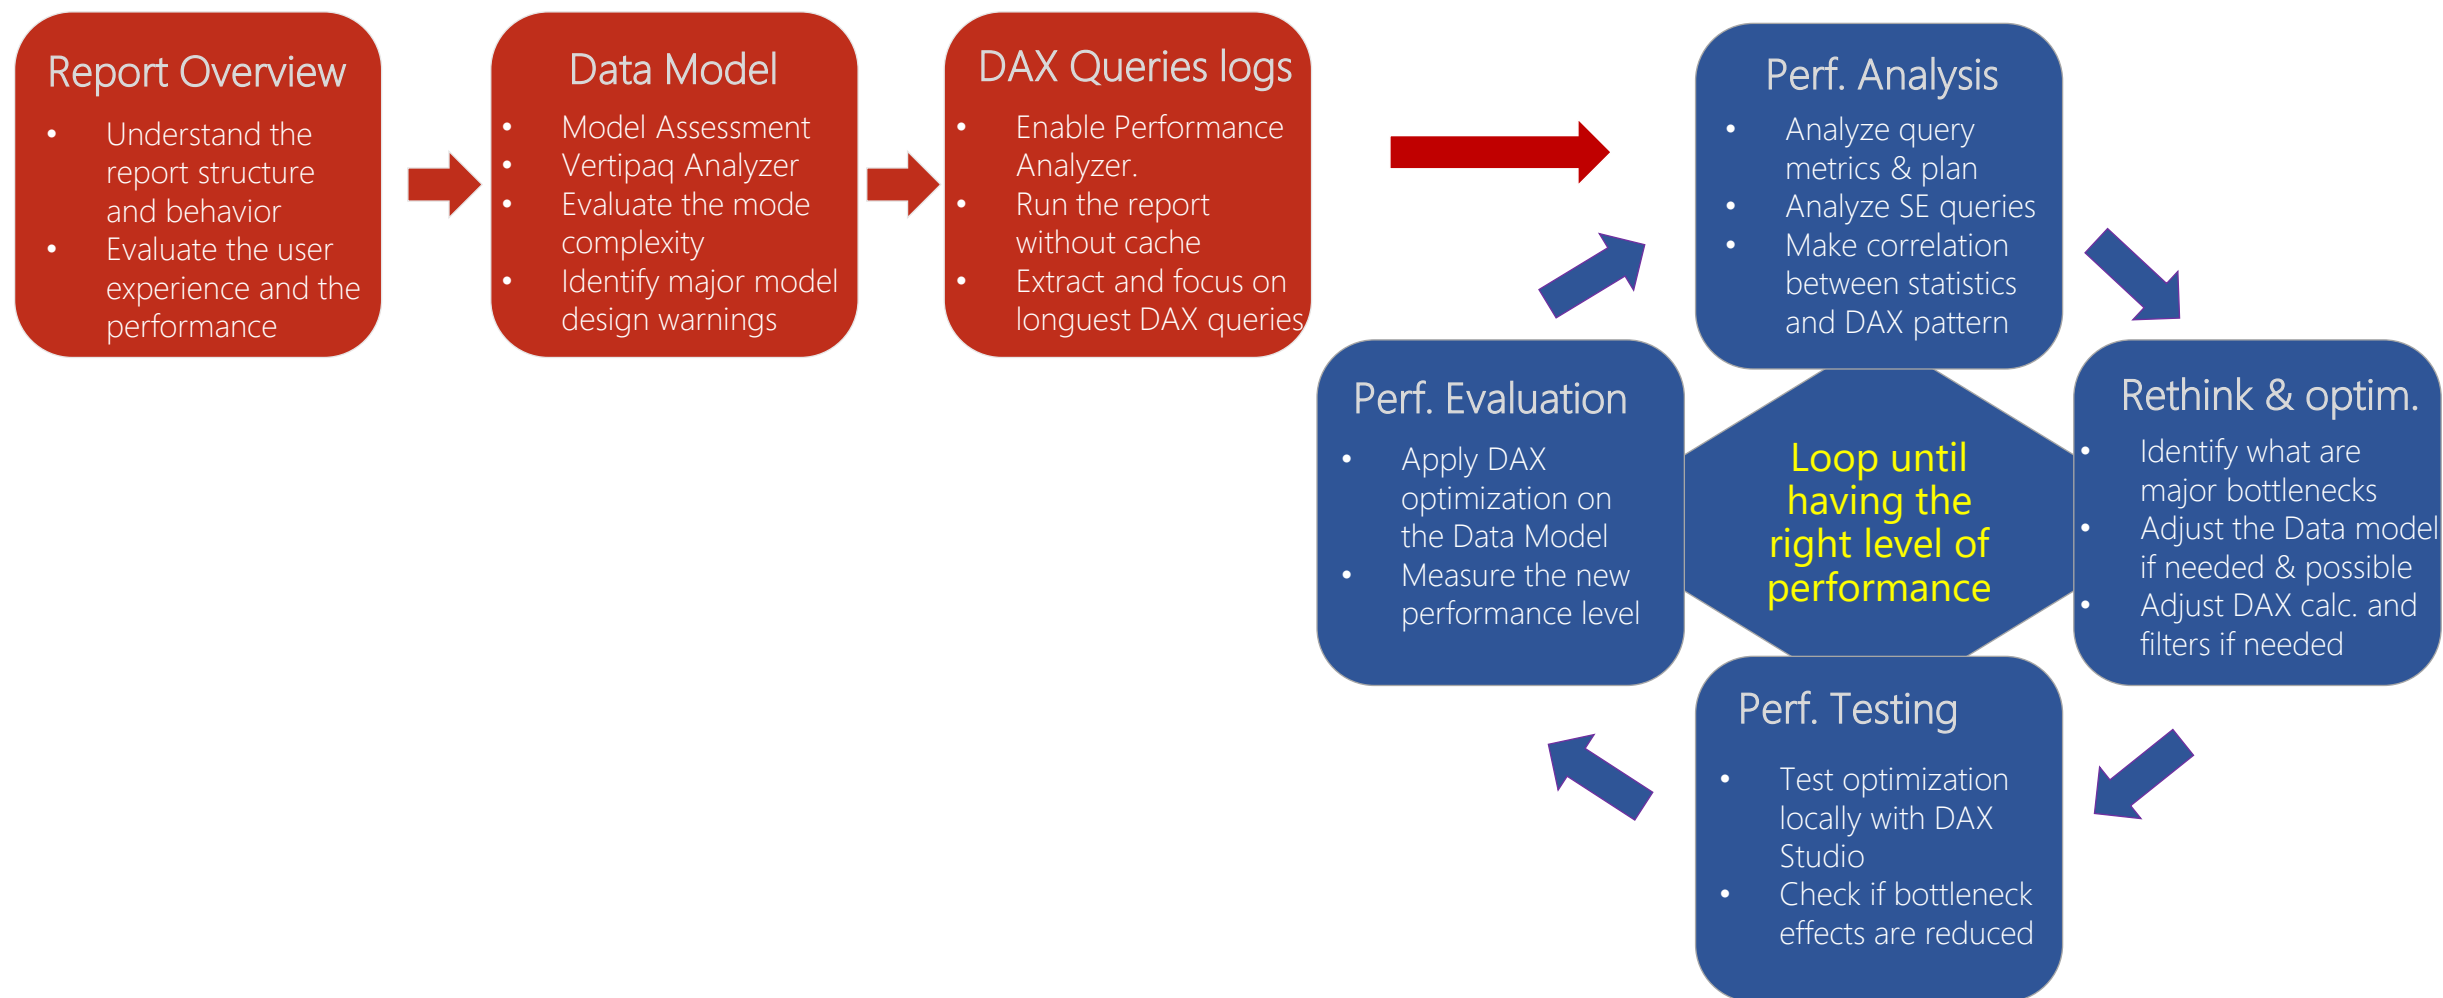
Clear Export

Name	Duration (ms)
Refreshed visual	-
Text box	23
Image	57
Text box	23
Slicer	101
Text box	24
Text box	28
Red - Card with States	158
DAX query	7
Visual display	54
Other	97
Copy query	
Run in DAX Query View	
Black - Card with States	159
Blue - Card with States	160

Report – Interaction

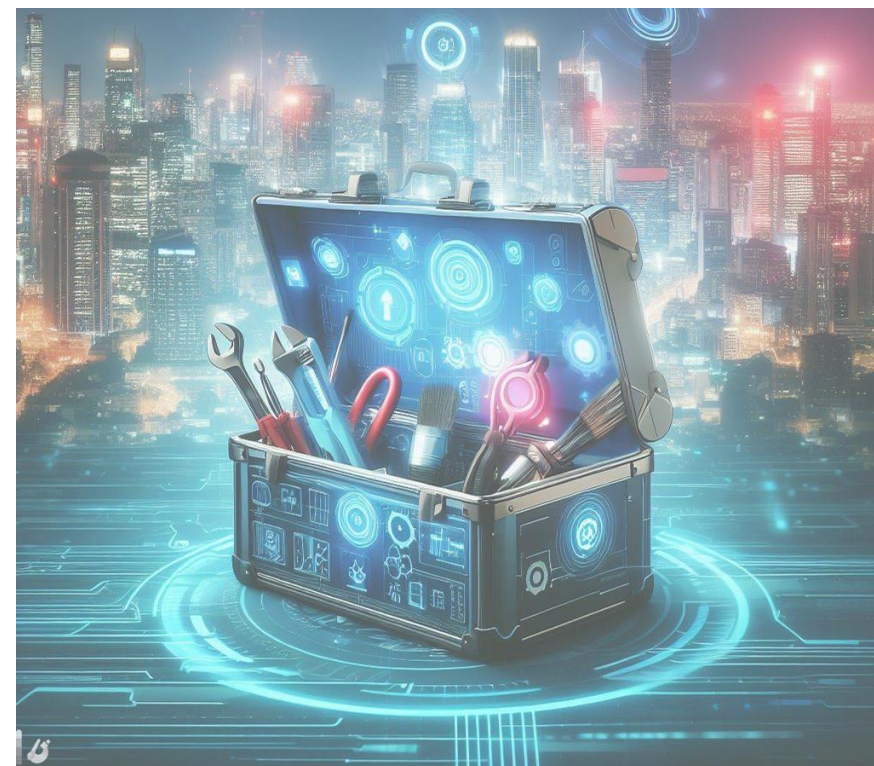
1. Limit the number of visuals in dashboards and reports
2. Remove unnecessary interactions between visuals
3. Reduce the amount of data loaded on page load
4. Use tooltips to provide more information on visuals and metrics
5. Allow users to personalize visuals in a report
6. Reduce queries
7. Use report backgrounds for static images
8. Avoid scrolls within the visual and on page
9. Test custom visual performance before use
10. Use drillthrough buttons instead of expecting users to right-click on data points

Performance Remediation Process

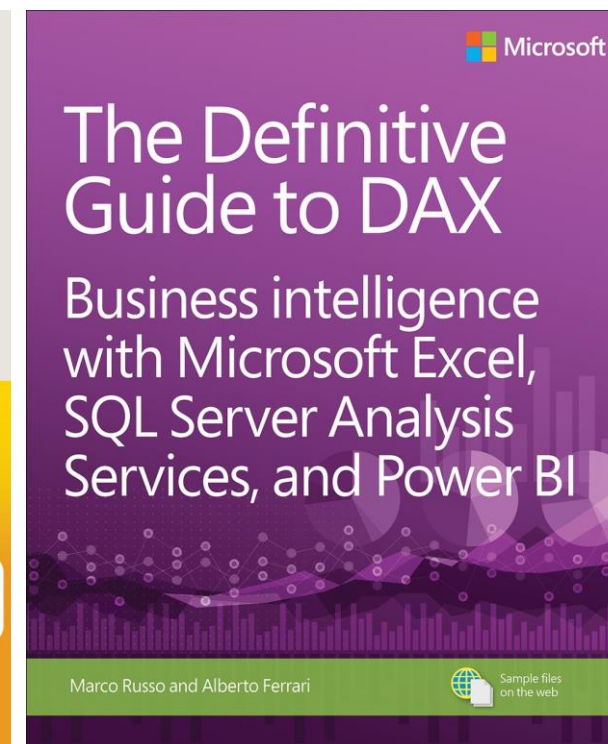
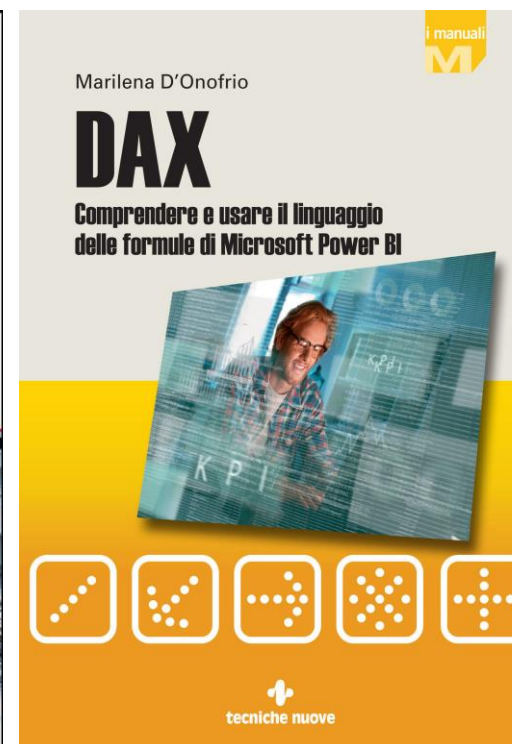
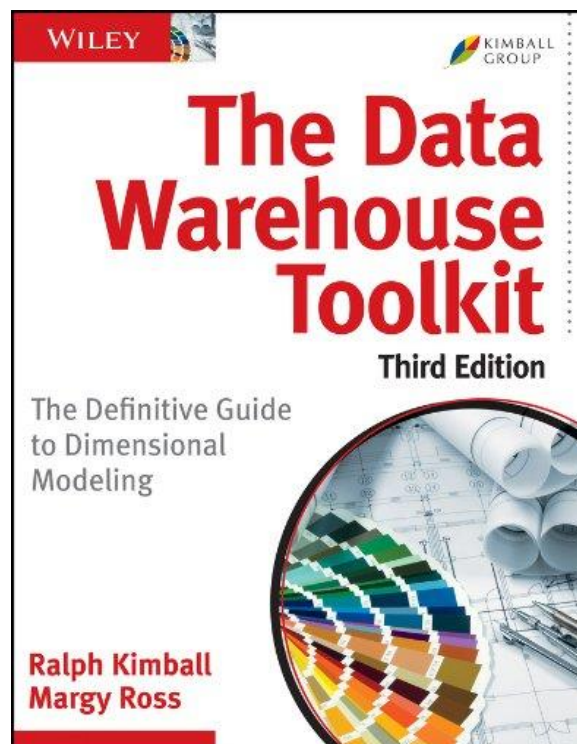
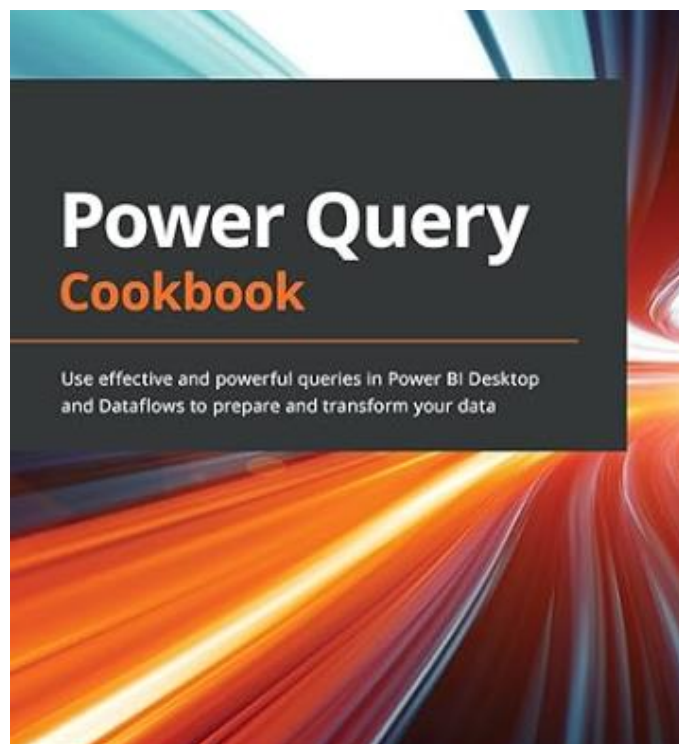


Tools

- <https://tabulareditor.com/downloads>
- <https://daxstudio.org/>
- <https://www.pbiexplorer.com/>
- <https://www.elegantbi.com/post/reportanalyzer>
- <https://powerbi.microsoft.com/en-my/blog/best-practice-rules-to-improve-your-models-performance/>
- <https://www.elegantbi.com/post/bestpracticерulesavings>



Book





28, 29, 30 NOVEMBRE NH MILANO CONGRESS CENTRE



Valuta la sessione
GRAZIE!